

Introduction to Diffusion Models

2022.01.03.

KAIST ALIN-LAB

Sangwoo Mo

Diffusion Model Boom!

- Diffusion model is SOTA on **image generation**
 - Beat BigGAN and StyleGAN on high-resolution images




Figure 1: Selected samples from our best ImageNet 512×512 model (FID 3.85)


Model	FID	sFID	Prec	Rec
LSUN Bedrooms 256×256				
DCTransformer [†] [42]	6.40	6.66	0.44	0.56
DDPM [25]	4.89	9.07	0.60	0.45
IDDPM [43]	4.24	8.21	0.62	0.46
StyleGAN [27]	2.35	6.62	0.59	0.48
ADM (dropout)	1.90	5.59	0.66	0.51
ImageNet 512×512				
BigGAN-deep [5]	8.43	8.13	0.88	0.29
ADM	23.24	10.19	0.73	0.60
ADM-G (25 steps)	8.41	9.67	0.83	0.47
ADM-G	7.72	6.57	0.87	0.42

Diffusion Model Boom!

- Diffusion model is SOTA on **density estimation**
 - Beat autoregressive models on likelihood score




(a) CIFAR-10 without data augmentation



(b) ImageNet 64x64

Diffusion Model Boom!

- Diffusion model is useful for **image editing**
 - Editing = Rough scribble + diffusion (i.e., naturalization)
 - Scribbled images are unseen for GANs, but diffusion models still can *denoise* them



Diffusion Model Boom!

- Diffusion model is useful for **image editing**
 - Also can be combined with vision-and-language model



“zebras roaming in the field”



“a girl hugging a corgi on a pedestal”



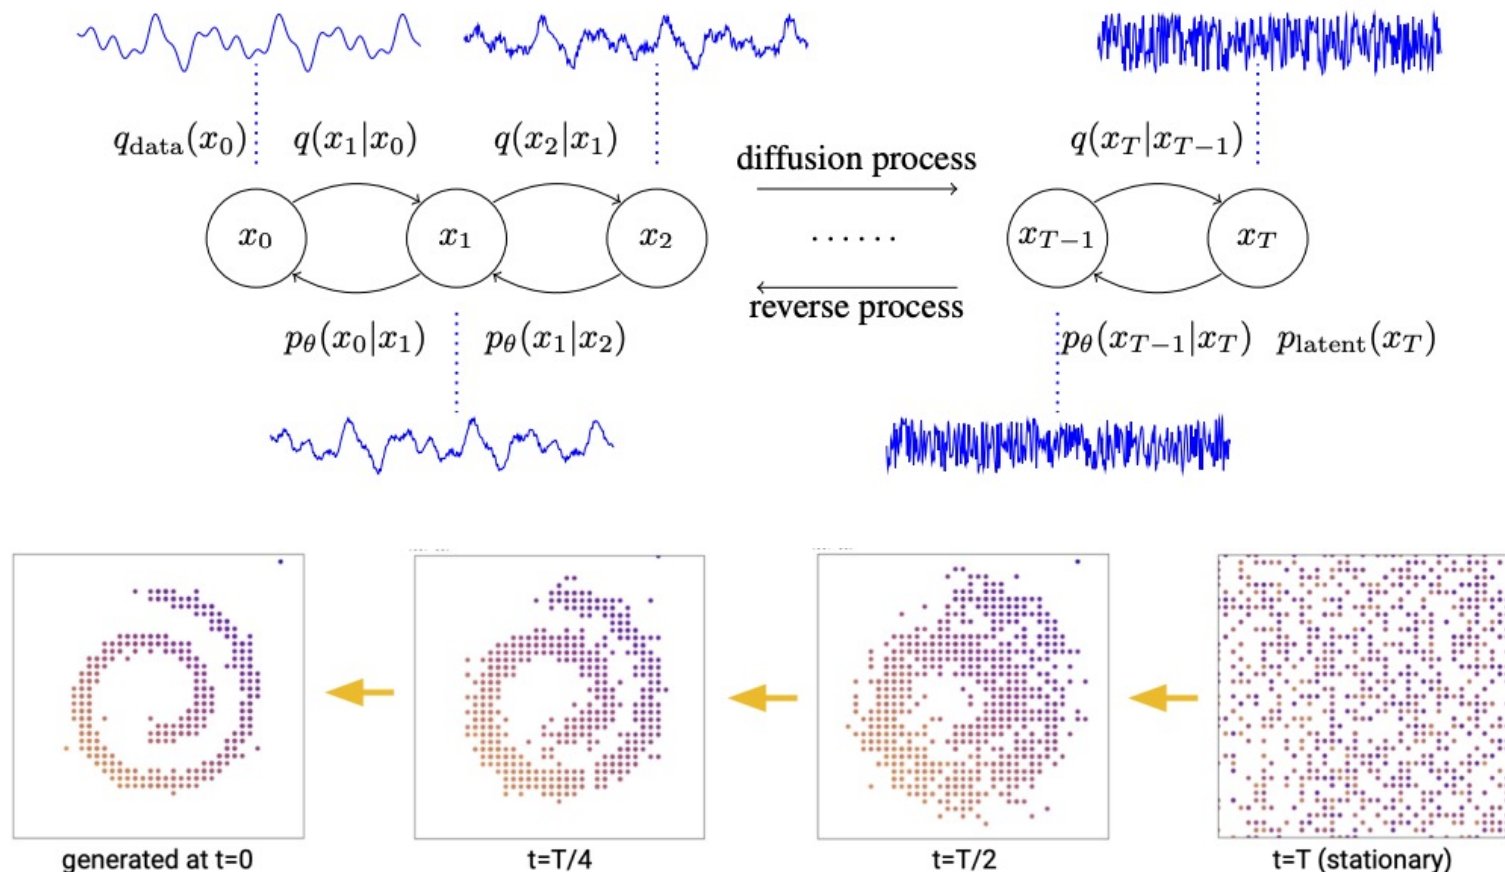
“a man with red hair”



“a vase of flowers”

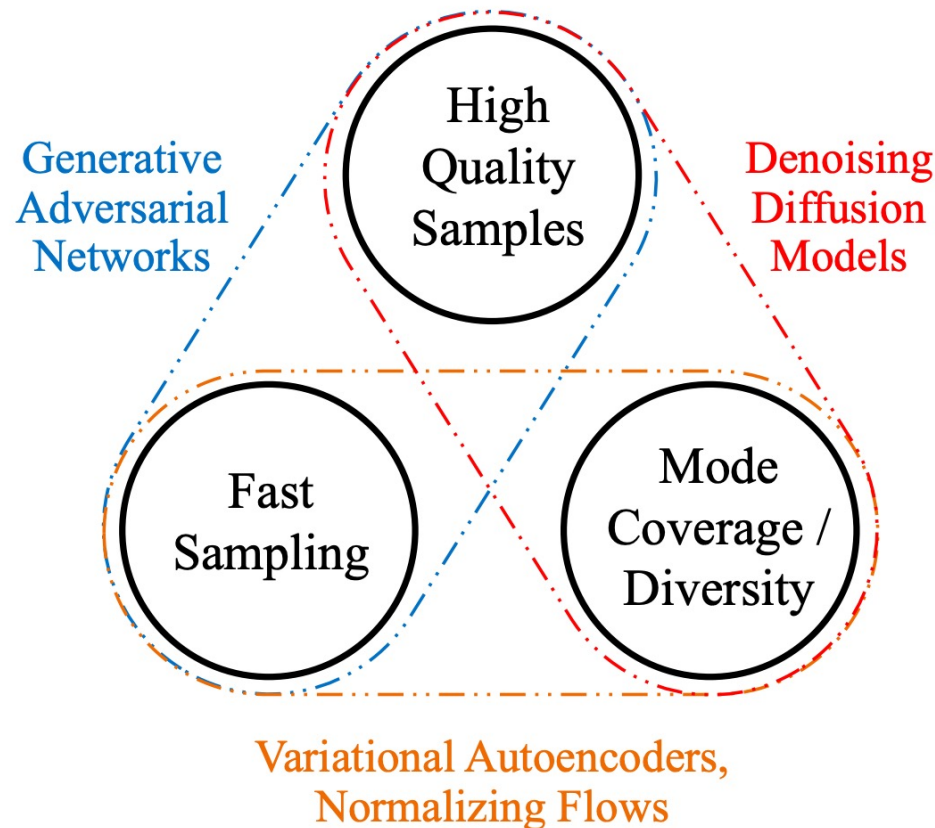
Diffusion Model Boom!

- Diffusion model is also effective for **non-visual** domains
 - Continuous domains like **speech**, and even for discrete domains like **text**



Diffusion Model is All We Need?

- **Trilemma** of generative models: Quality vs. Diversity vs. Speed
 - Diffusion model produces **diverse** and **high-quality** samples, but generations is **slow**



Outline

- **Today's content**


- Diffusion Probabilistic Model – ICML'15
- Denoising Diffusion Probabilistic Model (DDPM) – NeurIPS'20
 - Improve quality & diversity of diffusion model
- Denoising Diffusion Implicit Model (DDIM) – ICLR'21
 - Improve generation speed of diffusion model

- **Not covering**

- Relation of diffusion model and score matching
 - Extension to stochastic differential equation
 - There are lots of new interesting works (see NeurIPS'21, ICLR'22)
- See **Score SDE** (ICLR'21)

Diffusion Probabilistic Model

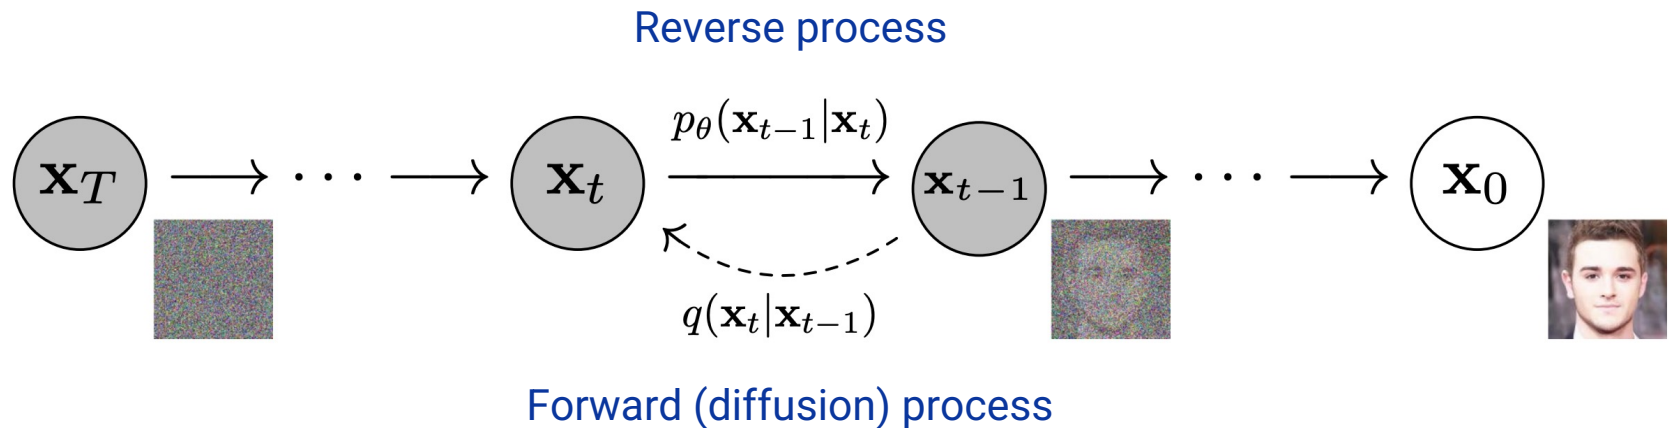
- Diffusion model aims to learn the **reverse of noise generation** procedure
 - **Forward step:** (Iteratively) Add noise to the original sample
 - The sample x_0 converges to the **complete noise** x_T (e.g., $\sim \mathcal{N}(0, I)$)



Forward (diffusion) process

Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure
 - **Forward step:** (Iteratively) Add noise to the original sample
 - The sample x_0 converges to the **complete noise** x_T (e.g., $\sim \mathcal{N}(0, I)$)
 - **Reverse step:** Recover the original sample from the noise
 - Note that it is the **“generation”** procedure



Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure
 - **Forward step:** (Iteratively) Add noise to the original sample
 - Technically, it is a product of **conditional noise** distributions $q(\mathbf{x}_t|\mathbf{x}_{t-1})$
 - Usually, the parameters β_t are **fixed** (one can jointly learn, but not beneficial)
 - **Noise annealing** (i.e., reducing noise scale $\beta_t < \beta_{t-1}$) is crucial to the performance

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}), \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) := \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$$

Diffusion Probabilistic Model

- **Diffusion model aims to learn the reverse of noise generation procedure**

- **Forward step:** (Iteratively) Add noise to the original sample

→ Technically, it is a product of **conditional noise** distributions $q(\mathbf{x}_t|\mathbf{x}_{t-1})$

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}), \quad q(\mathbf{x}_t|\mathbf{x}_{t-1}) := \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$$

- **Reverse step:** Recover the original sample from the noise

→ It is also a product of **conditional (de)noise** distributions $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$

- Use the **learned** parameters: **denoiser** $\boldsymbol{\mu}_\theta$ (main part) and **randomness** $\boldsymbol{\Sigma}_\theta$

$$p_\theta(\mathbf{x}_{0:T}) := p(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t), \quad p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t))$$

Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure

- **Forward step:** (Iteratively) Add noise to the original sample

Reverse step: Recover the original sample from the noise

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}), \quad p_{\theta}(\mathbf{x}_{0:T}) := p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t),$$

- **Training:** Minimize **variational lower bound** of the model $p_{\theta}(\mathbf{x}_0)$

$$\mathbb{E}[-\log p_{\theta}(\mathbf{x}_0)] \leq \mathbb{E}_q \left[-\log \frac{p_{\theta}(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right]$$

Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure

- **Forward step:** (Iteratively) Add noise to the original sample

Reverse step: Recover the original sample from the noise

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}), \quad p_{\theta}(\mathbf{x}_{0:T}) := p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t),$$

- **Training:** Minimize **variational lower bound** of the model $p_{\theta}(\mathbf{x}_0)$

$$\mathbb{E}[-\log p_{\theta}(\mathbf{x}_0)] \leq \mathbb{E}_q \left[-\log \frac{p_{\theta}(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right]$$

→ It can be decomposed to the **step-wise** losses (for each step t)

$$\mathbb{E}_q \left[\underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{L_T} + \sum_{t>1} \underbrace{D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))}_{L_{t-1}} \underbrace{-\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)}_{L_0} \right]$$

Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure

- **Training:** Minimize **variational lower bound** of the model $p_\theta(\mathbf{x}_0)$
→ It can be decomposed to the **step-wise** losses (for each step t)

$$\mathbb{E}_q \left[\underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{L_T} + \sum_{t>1} \underbrace{D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))}_{L_{t-1}} \underbrace{- \log p_\theta(\mathbf{x}_0|\mathbf{x}_1)}_{L_0} \right]$$

- Here, the **true reverse step** $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ can be computed as a **closed form** of β_t
 - Note that we only define the true forward step $q(\mathbf{x}_t|\mathbf{x}_{t-1})$


$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t^3 \mathbf{I})$$

$$\text{where } \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \tilde{\beta}_t^1 \mathbf{x}_0 + \tilde{\beta}_t^2 \mathbf{x}_t$$

- Since all distributions above are Gaussian, the KL divergences are tractable

Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure
 - **Network:** Use the **image-to-image translation** (e.g., U-Net) architectures
 - Recall that input is \mathbf{x}_t and output is \mathbf{x}_{t-1} , both are images
 - It is expensive since both input and output are high-dimensional
 - Note that the denoiser $\mu_\theta(\mathbf{x}_t, t)$ shares weights, but conditioned by step t




Diffusion Probabilistic Model

- Diffusion model aims to learn the **reverse of noise generation** procedure
 - **Sampling:** Draw a random noise x_T then apply the reverse step $p_\theta(\mathbf{x}_{t=1}|\mathbf{x}_t)$
 - It often requires the hundreds of reverse steps (very slow)



- Early and late steps change the high- and low-level attributes, respectively



Denoising Diffusion Probabilistic Model (DDPM)

- DDPM **reparametrizes** the reverse distributions of diffusion models
 - **Key idea:** The original reverse step **fully creates** the denoiser $\mu_{\theta}(\mathbf{x}_t, t)$ from \mathbf{x}_t
 - However, \mathbf{x}_{t-1} and \mathbf{x}_t share most information, and thus it is redundant
→ Instead, create the **residual** $\epsilon_{\theta}(\mathbf{x}_t, t)$ and add to the original \mathbf{x}_t

Denoising Diffusion Probabilistic Model (DDPM)

- DDPM **reparametrizes** the reverse distributions of diffusion models
 - **Key idea:** The original reverse step **fully creates** the denoiser $\mu_\theta(\mathbf{x}_t, t)$ from \mathbf{x}_t
 - However, \mathbf{x}_{t-1} and \mathbf{x}_t share most information, and thus it is redundant
→ Instead, create the **residual** $\epsilon_\theta(\mathbf{x}_t, t)$ and add to the original \mathbf{x}_t
 - Formally, DDPM **reparametrizes** the learned reverse distribution as¹

$$\mu_\theta(\mathbf{x}_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right)$$

and the step-wise objective L_{t-1} can be reformulated as²

$$\mathbb{E}_{t, \mathbf{x}_0, \epsilon} \left[\left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$$


1. α_t are some constants determined by β_t

2. Note that we need no “intermediate” samples, and only compare the forward noise ϵ and reverse noise ϵ_θ conditioned on \mathbf{x}_0

Ho et al. Denoising Diffusion Probabilistic Models. NeurIPS'20

Denoising Diffusion Probabilistic Model (DDPM)

- DDPM **initiated** the diffusion model boom
 - Achieved SOTA on CIFAR-10, with high-resolution scalability
 - It produces more diverse samples than GAN (no mode collapse)



Denoising Diffusion Implicit Model (DDIM)

- DDIM **roughly sketches** the final sample, then **refine** it with the reverse process
 - **Motivation:**
 - Diffusion model is slow due to the **iterative procedure**
 - GAN/VAE creates the sample by **one-shot** forward operation
 - \Rightarrow Can we combine the advantages for **fast sampling** of diffusion models?
 - **Technical spoiler:**
 - Instead of naïvely applying diffusion model upon GAN/VAE, DDIM proposes a **principled approach** of rough sketch + refinement

Denoising Diffusion Implicit Model (DDIM)

- DDIM **roughly sketches** the final sample, then **refine** it with the reverse process
- Key idea:
 - Given \mathbf{x}_t , generate the **rough sketch** \mathbf{x}_0 and **refine** $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ ¹
 - Unlike original diffusion model, it is not a Markovian structure





Figure 1: Graphical models for diffusion (left) and non-Markovian (right) inference models.

1. Recall that the original diffusion model uses $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$
Song et al. Denoising Diffusion Implicit Models. ICLR'21

Denoising Diffusion Implicit Model (DDIM)

- DDIM **roughly sketches** the final sample, then **refine** it with the reverse process
 - Key idea:** Given \mathbf{x}_t , generate the **rough sketch** \mathbf{x}_0 and **refine** $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$



- Formulation:** Define the forward distribution $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ as


$$q_\sigma(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\sqrt{\alpha_{t-1}}\mathbf{x}_0 + \sqrt{1 - \alpha_{t-1} - \sigma_t^2} \cdot \frac{\mathbf{x}_t - \sqrt{\alpha_t}\mathbf{x}_0}{\sqrt{1 - \alpha_t}}, \sigma_t^2 \mathbf{I}\right)$$

then, the **forward process** is derived from Bayes' rule

$$q_\sigma(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) = \frac{q_\sigma(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q_\sigma(\mathbf{x}_t|\mathbf{x}_0)}{q_\sigma(\mathbf{x}_{t-1}|\mathbf{x}_0)}$$

Denoising Diffusion Implicit Model (DDIM)

- DDIM **roughly sketches** the final sample, then **refine** it with the reverse process
 - Key idea:** Given \mathbf{x}_t , generate the **rough sketch** \mathbf{x}_0 and **refine** $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$




- Formulation:** Forward process is $q_\sigma(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) = \frac{q_\sigma(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q_\sigma(\mathbf{x}_t|\mathbf{x}_0)}{q_\sigma(\mathbf{x}_{t-1}|\mathbf{x}_0)}$

and reverse process is

$$\mathbf{x}_{t-1} = \underbrace{\sqrt{\alpha_{t-1}} \left(\frac{\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_\theta^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}} \right)}_{\text{"predicted } \mathbf{x}_0"} + \underbrace{\sqrt{1 - \alpha_{t-1} - \sigma_t^2} \cdot \epsilon_\theta^{(t)}(\mathbf{x}_t)}_{\text{"direction pointing to } \mathbf{x}_t"} + \underbrace{\sigma_t \epsilon_t}_{\text{random noise}}$$

Denoising Diffusion Implicit Model (DDIM)

- DDIM **roughly sketches** the final sample, then **refine** it with the reverse process
 - Key idea:** Given \mathbf{x}_t , generate the **rough sketch** \mathbf{x}_0 and **refine** $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$



- Formulation:** Forward process is $q_{\sigma}(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) = \frac{q_{\sigma}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)q_{\sigma}(\mathbf{x}_t|\mathbf{x}_0)}{q_{\sigma}(\mathbf{x}_{t-1}|\mathbf{x}_0)}$


and reverse process is
$$\mathbf{x}_{t-1} = \underbrace{\sqrt{\alpha_{t-1}} \left(\frac{\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}} \right)}_{\text{"predicted } \mathbf{x}_0"} + \underbrace{\sqrt{1 - \alpha_{t-1} - \sigma_t^2} \cdot \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}_{\text{"direction pointing to } \mathbf{x}_t"} + \underbrace{\sigma_t \epsilon_t}_{\text{random noise}}$$

- Training:** The variational lower bound of DDIM is **identical** to the one of DDPM¹
 - It is surprising since the forward/reverse formulation is totally different

1. Precisely, the bound is different, but the solution is identical under some assumption (though violated in practice)
Song et al. Denoising Diffusion Implicit Models. ICLR'21

Denoising Diffusion Implicit Model (DDIM)

- DDIM significantly reduces the **sampling steps** of diffusion model
 - Creates the outline of the sample after only 10 steps (DDPM needs hundreds)



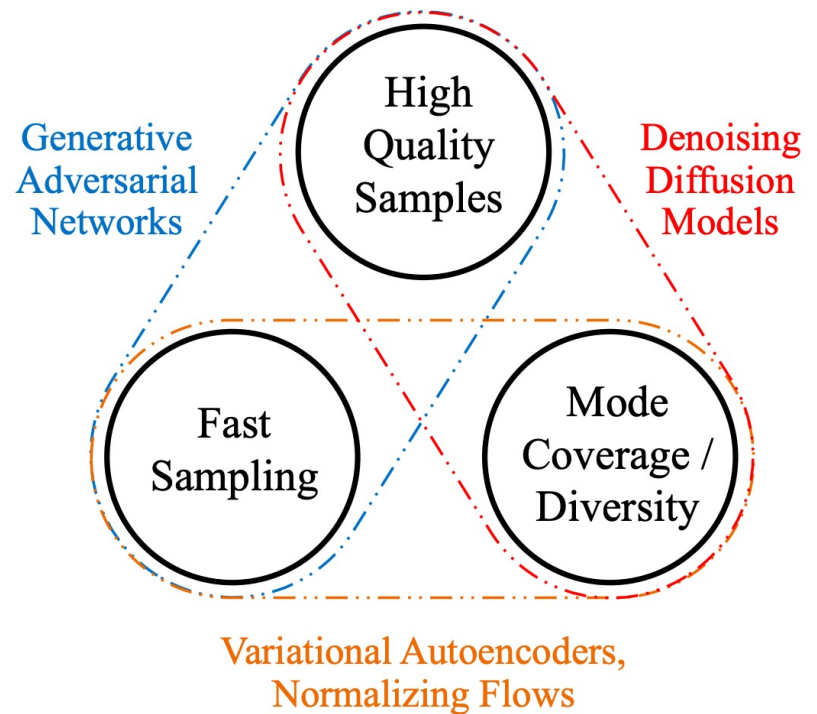
Take-home Message

- **New golden era of generative models**

- Competition of various approaches: GAN, VAE, flow, diffusion model¹
- Also, lots of hybrid approaches (e.g., score SDE = diffusion + continuous flow)

- **Which model to use?**

- **Diffusion model** seems to be a nice option for **high-quality** generation
- However, **GAN** is (currently) still a more practical solution which needs **fast sampling** (e.g., real-time apps.)



1. VAE also shows promising generation performance (see NVAE, very deep VAE)

Thank you for listening! 😊